

1

## Review: Minimizing Squared Error

- For a chosen set of weights,  $\mathbf{w}$ , used in linear regression, we define error as the (squared) difference between hypothesis function and correct output, summed over all test-cases:

$$Loss(\mathbf{w}) = \sum_{j=1}^N (y_j - h_{\mathbf{w}}(\mathbf{x}_j))^2$$

- Learning is then the process of finding a weight-sequence that **minimizes** this loss:


$$\mathbf{w}^* = \arg \min_{\mathbf{w}} Loss(\mathbf{w})$$

2

## Gradient Descent


$$Loss(\mathbf{w}) = \sum_{j=1}^N (y_j - h_{\mathbf{w}}(\mathbf{x}_j))^2$$

- The loss function forms a **contour** (here shown for one-dimensional data)
- For any initial set of weights ( $\mathbf{w}_0$ ) we are at some point on this contour



3


## Gradient Descent



- At this point, the derivate of the loss function points “uphill”
- The gradient descent update moves along the function in the **opposite** direction, to decrease loss most significantly

4

## Gradient Descent



- ▶ At **other** points, the derivative points “downhill” **already**
  - ▶ Gradient descent thus moves along the function in the **same** direction to decrease loss

Machine Learning (COMP 135)

5

## The Loss Gradient

$$\mathcal{L} = \sum_{j=1}^n (y_j - h_{\mathbf{w}}(\mathbf{x}_j))^2 = \sum_{j=1}^n (y_j - \mathbf{w} \cdot \mathbf{x}_j)^2$$

- ▶ For this loss function, the gradient with respect to any single weight is the first derivative of the loss applied to that weight:

$$\nabla \mathcal{L} = \frac{\partial \mathcal{L}}{\partial w_i} = -2 \sum_{j=1}^n x_{j,i} (y_j - \mathbf{w} \cdot \mathbf{x}_j)$$

- ▶ We can then modify that weight by **subtracting** the gradient:
  1. If the gradient is **positive** along the weight-axis, we are **decreasing** the weight to move in the **opposite** direction
  2. If the gradient is **negative**, we are **increasing** the weight to move in the **same** direction

Machine Learning (COMP 135)

6

## Modifying the Weight Updates

- ▶ In theory, we could modify a weight by applying the gradient **directly**:

$$w_i \leftarrow (w_i - \frac{\partial \mathcal{L}}{\partial w_i}) = (w_i + 2 \sum_{j=1}^n x_{j,i} (y_j - \mathbf{w} \cdot \mathbf{x}_j))$$


- ▶ In practice, however, this does not work well
  - ▶ Changing weights too much can “over-shoot” minimal points in the loss gradient
- ▶ Instead, we apply a **step-size parameter** to the weights
  - ▶ A multiplier,  $\alpha$ , most often  $< 1$ , to decrease the magnitude of update

$$w_i \leftarrow (w_i - \alpha \frac{\partial \mathcal{L}}{\partial w_i}) = (w_i + 2\alpha \sum_{j=1}^n x_{j,i} (y_j - \mathbf{w} \cdot \mathbf{x}_j))$$

Machine Learning (COMP 135)

7

## Convergence of Gradient Descent



- ▶ In the presence of large changes to the weights, the result can “ping pong” around the loss space in a way that **never** settles near a minimum
- ▶ Also known as the **learning rate**,  $\alpha$  provides a control parameter for this process
  1. This can be **fixed** to some small constant:  $\alpha = 0.001$
  2. Or, we may **decay** the parameter, making it smaller over time, decreasing it as a function of  $t$ , the number of iterations of the process:  $\alpha_0 = C \quad \alpha_t = \frac{C}{t} \quad (t \geq 1)$

Machine Learning (COMP 135)

8

## A Mathematical Convenience

- ▶ We can use a trick to make gradient computation more efficient, multiplying our loss function by a constant,  $c$ :

$$\mathcal{L}' = c \times \mathcal{L} = c \sum_{j=1}^n (y_j - \mathbf{w} \cdot \mathbf{x}_j)^2$$

- ▶ For a **positive constant**, this **does not** affect the target, minimizing weights:

$$\arg \min_{\mathbf{w}} \mathcal{L}' = \arg \min_{\mathbf{w}} c \mathcal{L} = \arg \min_{\mathbf{w}} \mathcal{L}$$

- ▶ Furthermore, the gradient for a single weight is also simple:

$$\nabla \mathcal{L}' = \frac{\partial \mathcal{L}'}{\partial w_i} = -2c \sum_{j=1}^n x_{j,i} (y_j - \mathbf{w} \cdot \mathbf{x}_j)$$

9

## A Mathematical Convenience

$$\nabla \mathcal{L}' = \frac{\partial \mathcal{L}'}{\partial w_i} = -2c \sum_{j=1}^n x_{j,i} (y_j - \mathbf{w} \cdot \mathbf{x}_j)$$

- ▶ Given this form of the gradient, we can simplify things by setting our multiplier  $c = 1/2$ , which means that the derivative calculation reduces somewhat:

$$\mathcal{L}' = \frac{1}{2} \sum_{j=1}^n (y_j - \mathbf{w} \cdot \mathbf{x}_j)^2$$

$$\nabla \mathcal{L}' = \frac{\partial \mathcal{L}'}{\partial w_i} = - \sum_{j=1}^n x_{j,i} (y_j - \mathbf{w} \cdot \mathbf{x}_j)$$

10

## A Mathematical Convenience


$$\nabla \mathcal{L}' = \frac{\partial \mathcal{L}'}{\partial w_i} = - \sum_{j=1}^n x_{j,i} (y_j - \mathbf{w} \cdot \mathbf{x}_j)$$

- ▶ This final form of the gradient gives us the weight update rule for linear regression
  - ▶ Descending the gradient by **subtracting** it out updates each
  - ▶ We use the learning rate to control the **speed** of this descent (i.e., the exact amount we are subtracting)

$$w_i \leftarrow w_i + \alpha \sum_j x_{j,i} (y_j - h_{\mathbf{w}}(\mathbf{x}_j))$$

11


## Potential Issues in Gradient Descent



- ▶ The squared-error loss for linear regression has a **convex** functional form
  - ▶ This means that, handled properly, gradient descent will converge upon a solution that is (very close to) **optimal**

12

## Potential Issues in Gradient Descent



- ▶ When loss functions are complex and non-convex, descending the gradient **may not** guarantee optimality
  - ▶ **Local minima** in the loss function are possible
  - ▶ Can be dealt with by a variety of techniques, e.g. randomly repeating starts
  - ▶ Can often be tolerated, so long as a **reasonable** minimum is found
  - ▶ We will see such non-convex scenarios later in the course

Machine Learning (COMP 135) 13

13

## Variants of Gradient Descent: Batch

- ▶ The gradient descent technique we described updates weights across **all** data in the training set:

1. Loop over all weights  $w_i$ , updating them:

$$w_i \leftarrow w_i + \alpha \sum_j x_{j,i} (y_j - h_{\mathbf{w}}(\mathbf{x}_j))$$

2. Stop on convergence

- ▶ This is also known as **batch gradient descent**
  - ▶ A very **stable** algorithm: as long as learning rate  $\alpha$  is not too large, will converge well to optimal or near-optimal solutions
  - ▶ Can be quite slow when data-set is large

Machine Learning (COMP 135) 14

14

## Variants of Gradient Descent: Stochastic

- ▶ Another version is **stochastic** gradient descent, where we use only **one** data-point at a time:

1. Pick some data-point  $\mathbf{x}_j$
2. Loop over all weights  $w_i$ , updating them:

$$w_i \leftarrow w_i + \alpha (x_{j,i} (y_j - h_{\mathbf{w}}(\mathbf{x}_j)))$$

3. Stop on convergence

- ▶ A faster technique, but less stable
  - ▶ We must be careful to reduce  $\alpha$  slowly in order to converge to a good solution

Machine Learning (COMP 135) 15

15

## Variants of Gradient Descent: Mini-Batch

- ▶ An “in-between” version is **mini-batch gradient descent**
- ▶ Like the batch version, we sum the overall weighted error, but only do so over some **fixed-size proper subset** of the data at any point
  - ▶ Can be more efficient than batch when data-set is very large
  - ▶ More stable than stochastic method, but care still needed to ensure that the algorithm will converge properly to minima
- ▶ Each of these approaches is used with many ML models
  - ▶ All that effectively changes is the loss function used, its gradient, and the resulting derivative calculation

Machine Learning (COMP 135) 16

16