

1

Reminder: Threshold Functions

1. We have data-points with n features:
 $\mathbf{x} = (x_1, x_2, \dots, x_n)$
2. We have a linear function defined by $n+1$ weights:
 $\mathbf{w} = (w_0, w_1, w_2, \dots, w_n)$
3. We can write this linear function as:
 $\mathbf{w} \cdot \mathbf{x}$
4. We can then find the **linear boundary**, where:
 $\mathbf{w} \cdot \mathbf{x} = 0$
5. And use it to define our **threshold** between classes:

$$h_{\mathbf{w}} = \begin{cases} 1 & \mathbf{w} \cdot \mathbf{x} \geq 0 \\ 0 & \mathbf{w} \cdot \mathbf{x} < 0 \end{cases}$$

Outputs 1 and 0 here are **arbitrary labels** for one of two possible classes


Machine Learning (COMP 135) 2

2

Hard Thresholds are Hard!

$$h_{\mathbf{w}} = \begin{cases} 1 & \mathbf{w} \cdot \mathbf{x} \geq 0 \\ 0 & \mathbf{w} \cdot \mathbf{x} < 0 \end{cases}$$

- ▶ The hard threshold function used by the perceptron algorithm (among others) produces some conceptual and mathematical challenges
- ▶ Gives a yes/no answer everywhere, which can be tricky when our data isn't linearly separable



Function is discontinuous (non-differentiable) at $x = 0$


Machine Learning (COMP 135) 3

3

The Logistic Function

$$h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w} \cdot \mathbf{x}}}$$

- ▶ We can generate a smooth curve by instead using the **logistic** function as a threshold
- ▶ We can treat this value as a **probability** of belonging to one class or another



Probability function is 0.5 at $x = 0$


Machine Learning (COMP 135) 4

4

Using the Logistic for Classification

$$h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w} \cdot \mathbf{x}}}$$

- ▶ Treated as a probability, the logistic can still be used to *classify* data, where the class is the one that has highest probability overall, while also supplying a probability for that outcome




A "coin flip" where we have $x = 0$

Machine Learning (COMP 135)

5

5

Issues with Linear Classification



- ▲ heart attack
- no heart attack


- ▶ Consider data about heart-attack risk, based upon body mass index (BMI) and blood pressure (BP)
- ▶ Even assuming linearly separable training data, linear classification gives a hard cut-off that may not be appropriate

Machine Learning (COMP 135)

6

6

Issues with Linear Classification



- ▲ heart attack
- no heart attack
- don't know?


- ▶ Given that **multiple** possible lines can separate this data, how do we classify a new instance when it lies in the region between the training instances?

Machine Learning (COMP 135)

7

7

Issues with Linear Classification



- ▲ heart attack
- no heart attack
- don't know?


- ▶ Even if we did settle on some fixed line, what do we do with something that is *very close* to the separator?

Machine Learning (COMP 135)

8

8

Using Probabilistic Classification



▲ heart attack □ don't know?
● no heart attack


- Logistic regression also generates a linear separator (where the weight-function = 0), but now it is giving us an empirical **distribution** over the data-set
- A new data point close to the line still has **some positive probability** of being in the class on the other side of it

Machine Learning (COMP 135)

9

9

Properties of the Logistic Function



$$h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w} \cdot \mathbf{x}}}$$

- Also known as the **Sigmoid**, from the shape of its plot
- It always has a value in range:
 $0 \leq x \leq 1$
- The function is everywhere differentiable, and has a **derivative** that is easy to calculate, which turns out to be useful for learning:

$$h'_{\mathbf{w}}(\mathbf{x}) = h_{\mathbf{w}}(\mathbf{x})(1 - h_{\mathbf{w}}(\mathbf{x}))$$

Machine Learning (COMP 135)

10

10

Logistic Regression

- In perceptron learning we update the weight vector in each case based upon a mis-classified instance, using the equation:

$$w_j \leftarrow w_j + \alpha(y_i - h_{\mathbf{w}}(\mathbf{x}_i)) \times x_{i,j}$$

- For the logistic, using the same loss function (squared error), we would do the same, but add an extra term:

$$w_j \leftarrow w_j + \alpha(y_i - h_{\mathbf{w}}(\mathbf{x}_i)) \times \underbrace{h_{\mathbf{w}}(\mathbf{x}_i)}_{\text{The difference between what output should be, and what our weights make it}} \times \underbrace{(1 - h_{\mathbf{w}}(\mathbf{x}_i))}_{\text{The derivative of the logistic}} \times \underbrace{x_{i,j}}_{\text{The } j\text{th feature-value}}$$

The difference between what output **should** be, and what our weights make it

The derivative of the logistic

Machine Learning (COMP 135)

11

11

Gradient Descent for Logistic Regression, I

- We could then use the same approach as for linear classification, starting with some random (or uniform) weights and then:

1. Choose an input \mathbf{x}_i from our data set that is wrongly classified.

2. Update vector of weights, $\mathbf{w} = (w_0, w_1, w_2, \dots, w_n)$:

$$w_j \leftarrow w_j + \alpha(y_i - h_{\mathbf{w}}(\mathbf{x}_i)) \times h_{\mathbf{w}}(\mathbf{x}_i)(1 - h_{\mathbf{w}}(\mathbf{x}_i)) \times x_{i,j}$$

3. Repeat until weights no longer change; modify learning parameter α over time to guarantee this.

- Again, we make α smaller and smaller over time, and the algorithm converges as $\alpha \rightarrow 0$


Machine Learning (COMP 135)

12

12

A Problem: Local Minima

$$w_j \leftarrow w_j + \alpha(y_i - h_{\mathbf{w}}(\mathbf{x}_i)) \times h_{\mathbf{w}}(\mathbf{x}_i)(1 - h_{\mathbf{w}}(\mathbf{x}_i)) \times x_{i,j}$$



- ▶ While this sort of weight update **does** drive the error down, it has a flaw: squared-error loss for logistic regression is **not** convex
 - ▶ Gradient descent can get stuck in locally optimal solutions that aren't ideal
 - ▶ Solution: change the error function!

13

Gradient Descent for Logistic Regression, II

- ▶ We change the weight-update in our gradient descent approach:

1. Choose an input \mathbf{x}_i from our data set that is wrongly classified.
2. Update vector of weights, $\mathbf{w} = (w_0, w_1, w_2, \dots, w_n)$:

$$w_j \leftarrow w_j + \alpha x_{i,j} (y_i - h_{\mathbf{w}}(\mathbf{x}_i))$$

3. Repeat until weights no longer change; modify learning parameter α over time to guarantee this.

- ▶ Note: the update **looks** like the update for linear regression, but:

1. It only uses a **single** incorrect data-point, not the sum of **all** errors
 2. The hypothesis h is an application of the logistic function
- ▶ The reason we can do this update is that we **don't** use the squared-error loss, but use a **different** loss function: **logistic loss**

14

Gradient Descent for Logistic Regression

$$w_j \leftarrow w_j + \alpha x_{i,j} (y_i - h_{\mathbf{w}}(\mathbf{x}_i))$$

- ▶ The logistic update equation, via gradient descent, minimizes the **log-loss** (also known as **logistic loss** or **binary cross entropy**):

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N y_i \log p_i + (1 - y_i) \log(1 - p_i)$$


- ▶ For these purposes, we treat the output of the logistic as the probability we are interested in:

$$p_i \triangleq h_{\mathbf{w}}(\mathbf{x}_i)$$

- ▶ Over time, we drive the loss towards 0

15

Gradient Descent for Logistic Regression



$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N y_i \log p_i + (1 - y_i) \log(1 - p_i)$$

- ▶ The log-loss is a **convex** function
 - ▶ Like the losses used in linear regression and the perceptron
- ▶ This means that the gradient descent process (with suitable values of α) will **converge** upon a near-optimal solution


16

Logarithmic Loss vs. Thresholded Error

- For an individual data element, the log loss is an **upper bound** on a threshold-based (1/0) loss:

$$\mathcal{E}(h_{\mathbf{w}}(\mathbf{x}_i), y_i) = \begin{cases} 0 & \text{if } h_{\mathbf{w}}(\mathbf{x}_i) = y_i \\ 1 & \text{if } h_{\mathbf{w}}(\mathbf{x}_i) \neq y_i \end{cases}$$

$$\mathcal{L}(h_{\mathbf{w}}(\mathbf{x}_i), y_i) = -[y_i \log h_{\mathbf{w}}(\mathbf{x}_i) + (1 - y_i) \log(1 - h_{\mathbf{w}}(\mathbf{x}_i))]$$



- This graph assumes:

- True label is 1
- Threshold used is 0.5 (i.e., $h_{\mathbf{w}} = 1$ if probability assigned is $p \geq 0.5$)
- Log base 2 is used

Machine Learning (COMP 135) 17

17

Linear vs. Logistic Regression

Linear Regression	Logistic Regression
A value $x \in \mathbb{R}$	A value $0 \leq x \leq 1$
Output value of an arbitrary function	Probability of belonging to a certain class
Tries to find line that best fits to the data	Tries to find separator that best divides the classes

Machine Learning (COMP 135) 18

18

Linear vs. Logistic Regression in Mathematical Terms

Linear	
Loss function	$Loss(\mathbf{w}) = \sum_{j=1}^N (y_j - h_{\mathbf{w}}(\mathbf{x}_j))^2$
Weight-update equation	$w_i \leftarrow w_i + \alpha \sum_j x_{j,i} (y_j - h_{\mathbf{w}}(\mathbf{x}_j))$
Logistic	
Loss function	$-\frac{1}{N} \sum_{j=1}^N [y_j \log h_{\mathbf{w}}(\mathbf{x}_j) + (1 - y_j) \log(1 - h_{\mathbf{w}}(\mathbf{x}_j))]$
Weight-update equation	$w_i \leftarrow w_i + \alpha x_{j,i} (y_j - h_{\mathbf{w}}(\mathbf{x}_j))$

Machine Learning (COMP 135) 19

19

Another Approach: ADALINE classifiers

- Rather than a perceptron or logistic approach, what if we tried to use linear regression itself to build a classifier?
- For two classes, we could:
 - Label data using two class-labels, $y \in \{+1, -1\}$
 - Fit a linear regression to this data using squared loss (now measured as the difference between the linear value and the class-label, not some other real number) and the same weight-updates as before
 - Classify data based upon whether the resulting linear function is ≥ 0 (in which case it is assigned +1) or not (-1)
- This is known as a least-squares or ADALINE (Adaptive Linear Neuron) classifier

Machine Learning (COMP 135) 20

20

Treatment of Outliers in Data

