


1

Neural Networks for Images


- ▶ A regular feed forward network can sometimes prove problematic for image-processing tasks
 - ▶ Given a (100×100) pixel color image, each with 3 color-channel (e.g. RGB) values, we end up with many, many weights to be learned
 - ▶ In addition, a 1-D weight-vector doesn't carry any real information about **spatial relationships** between image features (edges, blocks of color, ...)



2

Convolutional Neural Networks (CNNs)

- ▶ To capture image dynamics, and expand what the networks can do, we organize neurons into stacks of 3-dimensional volumes
 - ▶ Each is connected to later volumes, filtering and flattening down to the usual final $(C \times 1)$ classification-output layer (where C is the number of classes)



3


Types of Layers in CNNs

- ▶ **INPUT:** as in a typical NN, each neuron corresponds to a single **input feature-value**
 - ▶ Only the 3-D arrangement is different
- ▶ **OUTPUT:** again, as in a typical NN, these are **fully-connected** layers
 - ▶ Each neuron is connected to all of those in the volume above
 - ▶ Each computes a function, like the sigmoid (*softmax*), typically giving probabilities for each of the possible output classes
- ▶ **OTHER:** layers between can play different possible roles
 1. **CONVOLUTION:** transformations on sub-regions
 2. **RELU:** application of the $\max(0, x)$ function
 3. **POOLING:** down-sampling to reduce volume size

4

Deep Convolutional Networks

- For complex image-classification tasks, we may use many layers, combining the types in varying orders



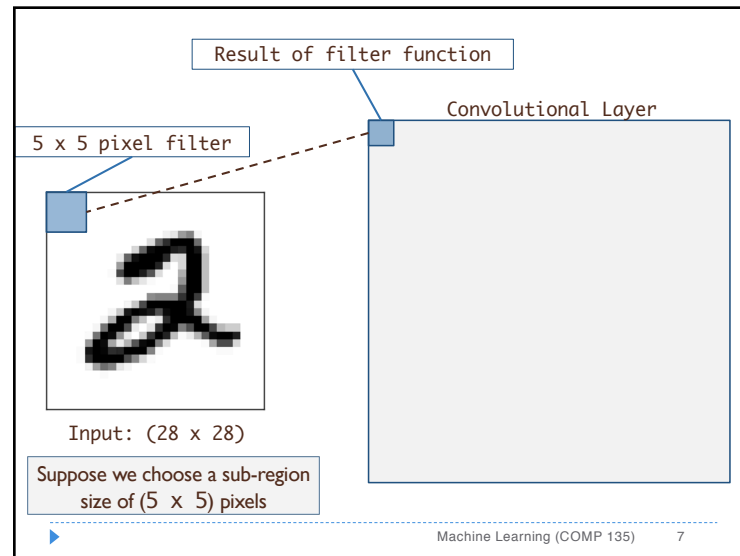
5

Convolutional (CONV) Layers

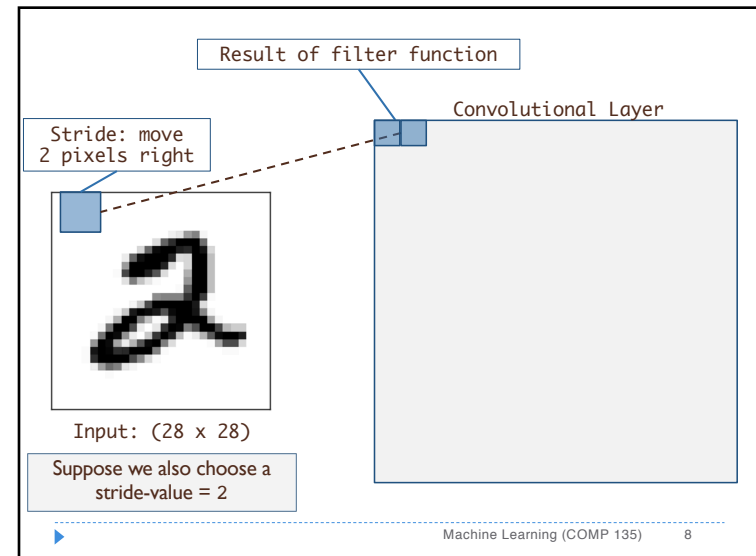
- The core innovation in a CNN is the idea of a **spatial filter**, which is a 3-D volume where:
 - Each neuron in one layer computes a function on a proper **sub-region** of the layer above
 - We form the CONV layer by “tiling” the prior layer, in (possibly) overlapping sub-regions
 - Every neuron in one layer shares a **single set** of weights, and so computes the same function
- Two main decisions in building such a layer:
 - What **size** of sub-region should we use?
 - What is our **stride**; i.e., **how far** do we move over each time we connect our next sub-region?

Machine Learning (COMP 135) 6

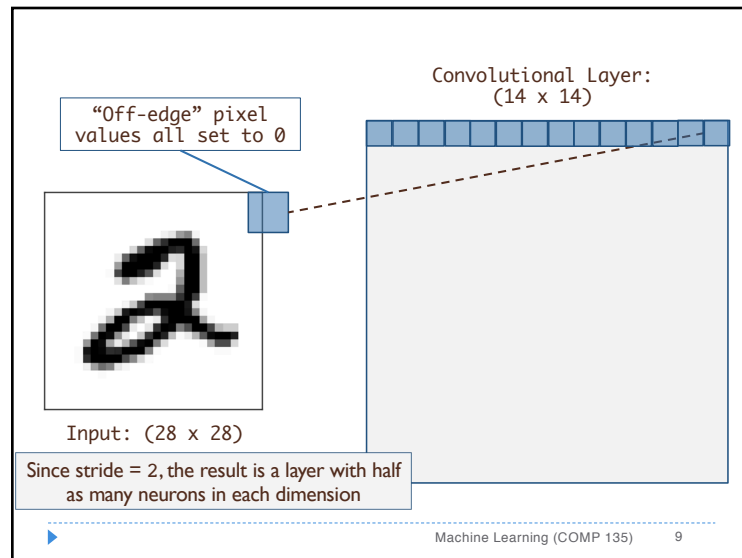
6



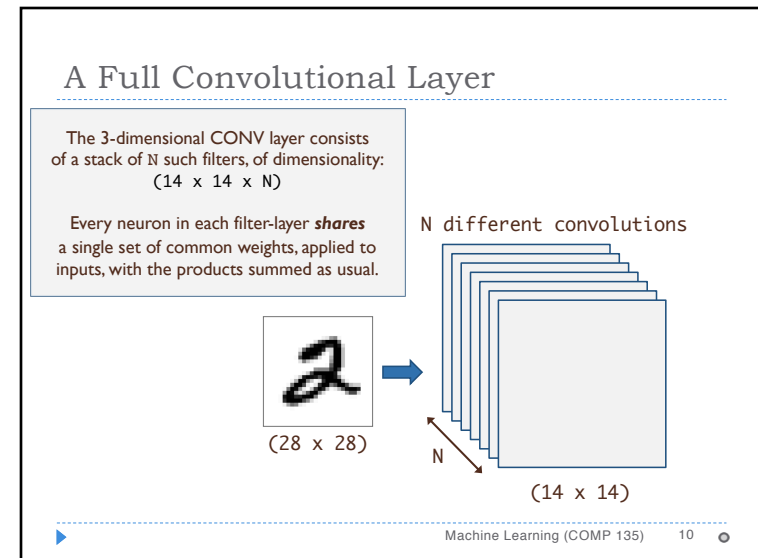
7



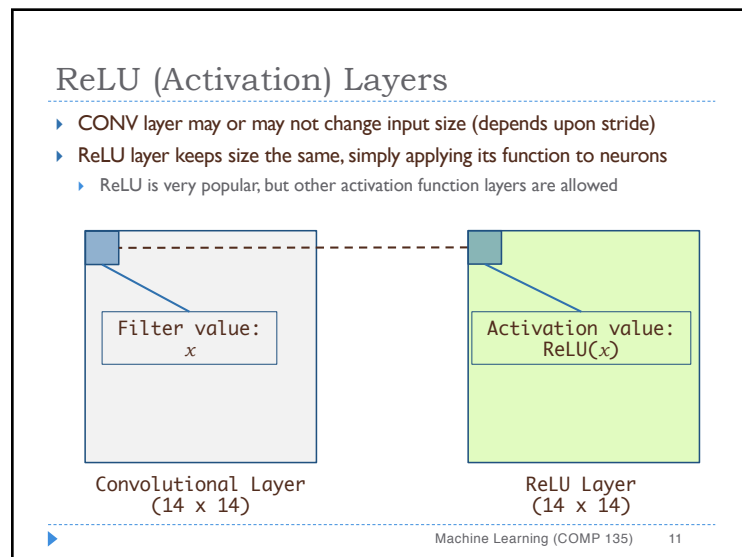
8



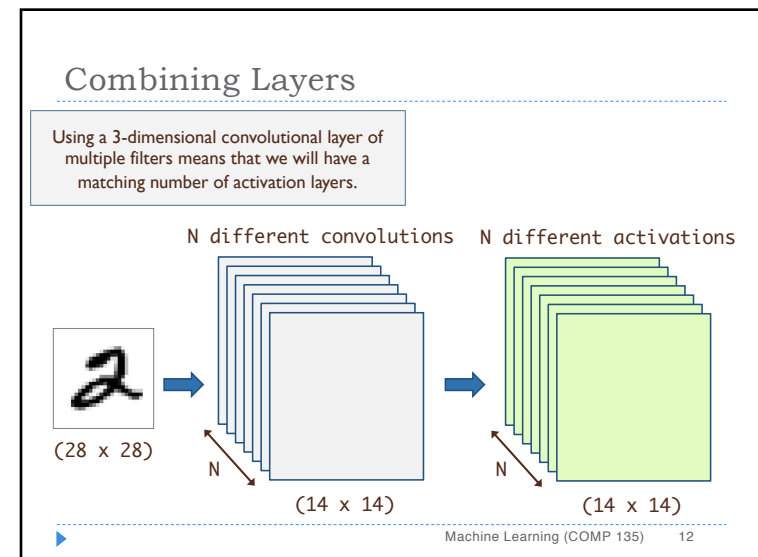
9



10




11



12

Pooling Layers

- While CONV and ReLU layers can compute more complex functions, POOL layers **down-sample** a region, reducing it to something simpler (usually its MAX value)




Machine Learning (COMP 135) 13

13

Pooling Layers

- Again, we stride across the layer, reducing the overall size by avoiding overlap
- Most common approach: (2 x 2) region, with stride = 2




Machine Learning (COMP 135) 14

14

Combining Layers

Again, each layer is 3-dimensional (until the final output layer).



Machine Learning (COMP 135) 15

15

Uses of CNNs and Other Deep Networks

- Convolutional networks have become increasingly popular for image and other spatial data
- Browser-based demos:
<https://cs.stanford.edu/people/karpathy/convnetjs/>
- A variety of applications of neural network models to a number of research problems
<https://youtu.be/Bui3DWs02h4>
<https://youtu.be/hPKJBXkyTKM>
<https://youtu.be/aKSILzbAqJs>
- Cat drawings!
<https://affinelayer.com/pixsrv/>

Machine Learning (COMP 135)

16

16